Using smartphone sensing and machine learning for personalized daily lapse risk prediction in a national sample of people with opioid use disorder

Kendra Wyant[aff-1] and John J. Curtin[aff-1]

[aff-1]Department of Psychology, University of Wisconsin-Madison

Abstract

Enter abstract here.

Keywords: Substance use disordersPrecision mental health

Introduction

About 4.5 million adults in the United States met criteria for opioid use disorder (OUD) in 2024 (**substanceabuseandmentalhealthservicesadministration2024NSDUHDetailed?**). OUD is associated with elevated mortality and morbidity, including increased risk of infectious disease, psychiatric comorbidity, and disability (**hserHighMortalityPatients2017?**). The societal burden of untreated OUD is substantial. The combined economic costs of lost productivity, health care expenditures, criminal justice involvement, loss of life, and reduced quality of life have been estimated to exceed $1 trillion (**florenceEconomicBurdenOpioid2021?**).

Medications for opioid use disorder (MOUD), including methadone and buprenorphine, are effective at reducing opioid use, craving, and overdose risk (**maEffectsMedicationassistedTreatment2019?**). They are also associated with improvements in mental health and social functioning, particularly when maintained over extended periods. However, despite strong evidence for MOUD, access remains limited and treatment dropout is common. Sustained retention in treatment is a key determinant of positive outcomes.

Lapses remain common even among individuals receiving MOUD and can increase the risk of treatment dropout, relapse, and overdose (**weinsteinLongtermRetentionOffice2017?**; **heinzLongitudinalObservationalStudy2025?**). Opioid lapses are particularly high-risk events. Reduced opioid tolerance (e.g., after a period of abstinence) and increased prevalence of fentanyl and other synthetic opioids in the drug supply substantially increases the likelihood of fatal overdose.

Identifying when an individual is at high lapse risk could enable earlier clinical intervention and potentially prevent relapse and overdose. However, risk for opioid use and lapse is dynamic and fluctuates across time in response to stress, craving, environmental cues, and treatment engagement. Smartphone sensing methods, including brief daily self-report surveys and passive geolocation sensing, offer a .

Methods

Transparency and Openness

We adhere to research transparency principles that are crucial for robust and replicable science. First, we published this study's protocol as a registered report (International Registered Report Identifier [IRRID]: DERR1-10.2196/29563) during the initial enrollment of pilot participants (Moshontz et al., 2021). Second, we followed the Transparent Reporting of a multivariable prediction model for Individual Prognosis Or Diagnosis extension for Artificial Intelligence (TRIPOD+AI) guidelines (Collins et al., 2024). A TRIPOD+AI checklist is available in the supplement. Finally, our features, labels, questionnaires, and other study materials are publicly available on our OSF page (https://osf.io/zvm7s/overview) and our annotated analysis scripts and results are publicly available on our study website (https://jjcurtin.github.io/study_risk2/).

Participants

We recruited participants in early recovery from opioid use disorder from across the United States. We recruited through national digital advertising and collaborations with treatment providers at MOUD clinics. Our recruitment strategy was designed to create a diverse sample

with respect to demographics (gender, age, race, and ethnicity), and geographic location (urban and rural). We required participants:

- were age 18 or older,
- could read, write, and speak in English,
- were enrolled in and adherent with an MOUD program for at least 1 month but no longer than 12 months or enrolled in or recently completed a day treatment program, and
- had an Android smartphone with an active data plan.

We did not exclude participants for comorbid substance use or other psychiatric disorders.

Procedure

Participants completed three video or phone visits over approximately 12 months. During the enrollment visit, study staff obtained written informed consent and collected demographic information. They walked participants through how to download the Smart Technology for Addiction Recovery (STAR) study app, provided a set of video tutorials to learn how to use the app, and instructed participants to complete the intake survey within the app. The STAR app was developed by the UW Madison Center for Health Enhancement Systems Studies and used for all data collection. Within the app participants could control their data sharing options, monitor completed study tasks, receive reminder notifications about tasks, message staff, and access STAR's suite of resources and tools for people in recovery from AUD (Gustafson et al., 2014). Enrolled participants met with study staff one week later to troubleshoot technical issues. At the end of the study enrollment period participants met briefly with study staff for a debriefing

session. While on study, participants were expected to complete daily surveys, monthly surveys, and share geolocation sensing data. Other sensing data streams (i.e., daily video check-ins, cellular communications, and app usage data) were collected as part of the parent grant's aims (NIDA R01 DA047315). All procedures were approved by the University of Wisconsin-Madison Institutional Review Board (Study #2019-0656).

Participants were compensated $40 for completing the enrollment visit and watching the STAR video tutorials and $10 for the two follow-up phone visits. Participants were compensated $10 per month for the monthly survey, up to $15 per month for the daily surveys[1], and $5 per month for sharing geolocation data.[2] Participants were also paid $50 per month to offset the cost of maintaining a cellphone data plan.

Measures

Individual Differences

We collected self-report information about demographics (age, gender, sexual orientation, race, ethnicity, education, employment, and income). Zip codes from participants' reported home addresses were linked to Rural–Urban Commuting Area (RUCA) codes to characterize the rural–urban status of their residences (Economic Research Service US Department of Agriculture, n.d.). These variables were collected to characterize our sample. They were also included as features in our model to allow for potential interactions with the sensing features and control for

[1] $0.25 per survey plus an additional $3.70 for completing at least 12 surveys in the first half of the month and an additional $3.70 for completing at least 12 surveys in the second half of the month (total range of $14.40 - $15.15)

[2] Participants were also compensated for providing other sensing data as part of the parent grant's aims, including $10 per month for sharing cellular communication logs and text message content and up to $10 per month for the daily video check-in.

demographic differences on lapse risk. Demographic variables with some evidence for influencing OUD treatment access and clinical outcomes and/or carrying general societal stigma (gender, race/ethnicity, income, sexual orientation, and geographic location) were used to perform subgroup analyses for evaluating model fairness (Greenfield et al., 2007; Kilaru et al., 2020; Lofaro et al., 2025; Martin et al., 2022; Olfson et al., 2022; Pinedo, 2019; Xin et al., 2023).

We collected information about OUD history to characterize OUD severity and recovery stability across our sample. These measures included self-reported Diagnostic and Statistical Manual of Mental Disorders, Fifth Edition OUD symptoms (American Psychiatric Association, 2022), past month opioid use, past month residential treatment for OUD, past month receipt of psychiatric medication, preferred opioid, preferred route of administration, and lifetime history of overdose. As part of the aims of the parent project, we collected many other trait and state measures throughout the study. A complete list of all measures can be found in our registered report (Moshontz et al., 2021).

Daily Survey

Participants completed a brief (1-2 minute) 16-item daily survey each day on study through the STAR app. The daily survey became available in the app at 5:00 AM CST each morning and participants had 24 hours to complete it. Participants could enable push notifications for reminder prompts to complete the survey. Participants reported opioid lapses on the first daily survey item. If participants responded “yes” to the question “Have you used any opioids for non-medical reasons that you have not yet reported?”, they were prompted to select the day(s) and time(s) of the lapse(s). Times were reported in 6-hour increments (12:00am–5:59am, 6:00am–11:59am,

12:00pm–5:59pm, 6:00pm–11:59pm). These reports served as the primary outcome for the lapse risk prediction model. On the remaining 15 items, participants reported any other drugs that they had used, whether they took their MOUD as prescribed in the past 24 hours, and rated the maximum intensity of recent (i.e., in the past 24 hours) experiences of pain, craving, risky situations, stressful events, and pleasant events. They also rated their sleep and how depressed, angry, anxious, relaxed, and happy they have felt in the past 24 hours. Lastly, participants responded to 2 future-facing items that asked about participants' motivation and confidence to continue to avoid using opioids for non-medical reasons over the next week. The full daily survey is available in the supplement.

Geolocation Sensing

The STAR app passively recorded participants' time-stamped geolocations (i.e., latitude and longitude) every 1.5-15 minutes, depending on their movement. We augmented the geolocation data with self-reported subjective context. On each monthly survey participants were asked 6 questions about frequently visited locations (i.e., locations that the participant spent more than 3 minutes on 2 or more times in a month) from the previous month. Participants were asked to describe the type of place, what they typically do there, the general frequency of pleasant and unpleasant experiences at the location, and the extent to which spending time there helps or harms their recovery. The full set of location questions is avilable in the supplement.

Data Analytic Strategy

Data preprocessing, modeling, and Bayesian analyses were done in R using the tidymodels ecosystem (Goodrich et al., 2023; Kuhn, 2022; Kuhn & Wickham, 2020). Models were trained and evaluated using high-throughput computing resources provided by the University of Wisconsin Center for High Throughput Computing (Center for High Throughput Computing, 2006).

Feature Engineering

Features were calculated using only data collected before the start of each prediction window to ensure our models were making true future predictions. We calculated a total of 297 features from three data sources:

1. *The prediction window day.* We created one-hot coded features for day of the week for the start of the prediction window.
2. *Demographic individual differences.* For demographics, we created ordinal features for age, education, and income based on predefined response ranges (see Table 1 for ranges), an ordinal feature for RUCA code associated with home address (range 1-10), and one-hot coded features for gender (male, not male), race and ethnicity (non-Hispanic White, Hispanic and/or not White), orientation (heterosexual, not heterosexual), and employment (employed, not employed/looking for work, not employed/not looking or able to work).
3. *Dynamic daily survey and geolocation sensing data.* For both sets of sensing features we calculated two types of features: raw and difference features. Raw features represent the feature value calculated within a given scoring epoch (e.g., the maximum urge rating reported

on the daily survey during the 168 hours immediately preceding the start of the prediction window). Difference features capture participant-level changes from their baseline scores. Specifically, we subtracted each participant's mean score for each feature using all available data prior to the prediction window from the associated raw feature (e.g., the participant's average urge rating across all prior daily surveys subtracted from the maximum urge rating in the preceding 168 hours).

For the daily survey, we calculated raw and difference features for the most recent and the min, max, and median responses in the 168 hours preceeding the start of the prediciton window for the 13 5-point likert scale items (items 4-16; pain, urge, risky situation, stressful event, pleasant event, sleep, depressed, angry, anxious, relaxed, happy, abstinence motivation, and abstinence confidence). We also calculated raw and difference rate features based on counts of previous opioid lapses, other drug use, missed MOUD doses (items 1-3), and completed daily surveys in 24 and 168-hour feature scoring epochs (e.g., counts of opioid lapses within the 168 hours preceding the start of a prediction window divided by 168).

We used two feature scoring epochs (24 and 168 hours before the start of the prediction window) to create proximal and distal geolocation features. Raw geolocation points were cross-checked against known locations with reported subjective context. We used a threshold of 50 meters for matching context to geolocation points. We calculated raw and difference features for sum duration of time out of the home in the evenings (i.e., not at their home between the hours of 7:00pm-4:00am). We calculated raw and difference features for sum duration of time spent at locations according to what they indicated they do at the

location (spend time with friends, socialize with new people, religious activities, relax, spend time with family, volunteer, receive mental health care, receive physical health care, receive MOUD treatment, drink alcohol, take classes, work), how pleasant and unpleasant their experiences typically are at the location (Always, Most of the time, Sometimes, Rarely, Never), and how helpful and harmful the location is to their recovery (Extremely, Considerably, Moderately, Mildly, Not at all). We also calculated a raw and difference feature of location variance (i.e., the extent to which a participant's location changes over a feature scoring epoch).

All features (sets 1-3 above) were included in our full model. We also fit three models with varying ablated sensing features from set 3. Specifically, these models ablated daily survey features, ablated geolocation features, and ablated daily survey and geolocation features combined to assess the incremental predictive value of these features.

Other feature engineering steps performed during cross-validation included imputing missing values (median imputation for numeric features and mode imputation for nominal features), dummy coding nominal features, normalizing features to have a mean of zero and standard deviation of 1, bringing outlying values (|z-score| > 5) to the fence, and removing zero and near-zero variance features as determined from held-in data. We selected coarse median/ mode methods for handling missing data due to the low rates of missing values and computational costs associated with more advanced forms of imputation (e.g., KNN imputation, multiple imputation). A sample feature engineering script (i.e., tidymodels recipe) containing all feature engineering steps is available on our OSF study page.

Lapse Labels

Prediction windows started at 6:00am in participants' local timezones and ended at 5:59am the next day. This window start time was selected to more closely approximate a typical wake-sleep cycle as opposed to midnight-to-midnight calendar day. For each participant the first prediction window began at 6:00am on their second day of participation and rolled forward day-by-day until their participation ended (i.e., the last prediction window ended at 5:59am on the day of their last recorded daily survey). Responses from the first daily survey item were used to label each prediction window as lapse or no lapse. For example, if a participant reported a lapse on December 12 from 12:00am-5:59am the prediction window spanning 6:00am on December 11 through 5:59am on December 12 was labeled as a lapse.

Cross-Validation

We used nested cross-validation for selection and evaluation of the final candidate model configurations. We used 1 repeat of 5-fold cross-validation on the inner loop for model selection and 6 repeats of 5-fold cross-validation on the outer loop for evaluation. Participants were grouped so that all of their data were always in the held-in or held-out fold to avoid bias introduced when predicting a participant's data from their own data. Folds were stratified so that all folds contained comparable proportions of individuals who lapsed while on study.

Model Configurations

We initially considered three statistical algorithms that differed in terms of assumptions and bias-variance tradeoff: elastic net, random forest, and XGBoost. Preliminary analyses using simple 6x5 cross-validation suggested that XGBoost was superior (see supplement for initial

performance estimates by algorithm). Additionally, we planned to use the Shapley additive explanations (SHAP) method for calculating feature importance which is optimized for XGBoost. Therefore we only considered XGBoost for the final nested cross-validation analyses.

Final candidate model configurations differed across sensible values for key hyperparameters and outcome resampling method (i.e., up-sampling and down-sampling of the outcome at ratios ranging from 5:1 to 1:1). All resampling was exclusively done in the held-in training data (i.e., held-out data were not resampled) to prevent biasing performance estimates (Vandewiele et al., 2021). All model configurations were fit for the full model and the three feature ablation models (ablated daily survey features, ablated geolocation features, and ablated daily survey and geolocation features combined).

Model Evaluation

We evaluated the best full model's probability predictions across three domains: discrimination, calibration, and overall performance. We follow recommendations for reporting measures and plots to characterize these performance domains (Van Calster et al., 2025).

auROC and Model Comparisons. Our performance metric for model selection and evaluation was the area under the receiver operating characteristic curve (auROC). auROC is an aggregate measure of discrimination across all possible decision thresholds. Specifically, it represents the probability that the model will assign a higher predicted probability to a randomly selected positive case (lapse) compared to a randomly selected negative case (no lapse).

We used a Bayesian hierarchical generalized linear model to estimate the posterior distribution and 95% credible intervals (CI) for auROC for the 30 held-out test sets (i.e., held-out

data in the outer loop of the nested cross-validation procedure) for our best full and ablated models. We used weakly informative, data-dependent priors to regularize and reduce overfitting. Priors were set as follows: residual standard deviation ~ normal(location=0, scale=exp(2)), intercept (after centering predictors) ~ normal(location=2.3, scale=1.3), the two coefficients for window width contrasts ~ normal (location=0, scale=2.69), and covariance ~ decov(regularization=1, concentration=1, shape=1, scale=1). We set two random intercepts to account for our resampling method: one for the repeat, and another for the fold nested within the repeat. We specified three model contrasts that compared the full model to each of the three ablated models (full vs. ablated daily survey features, full vs. ablated geolocation features, full vs. ablated daily survey and geolocation features combined). auROCs were transformed using the logit function and regressed as a function of model contrast. From the Bayesian model we obtained the posterior distribution for auROC for the full and baseline models. We reported the median posterior probability for auROC and 95% CIs for each model. We then conducted four Bayesian model comparisons to determine the probability that the full model differed systematically in performance compared to the three ablated models.

We performed five dichotomous subgroup analyses to assess the fairness of our model's predictions. Using the same 30 held-out test sets and the same modeling procedure as above, we calculated the median posterior probability and 95% CI for auROC for each model separately by gender (not male vs. male), race/ethnicity (Hispanic and/or non-White vs. non-Hispanic White), income (less than $25,000 vs. more than $25,000), sexual orientation (heterosexual vs. not heterosexual), and geographic location (rural vs. urban)^[We followed guidelines from the United

States Health Resources and Services Administration and define urban as an area where the primary commuting flow is within a metropolitan core of 50,000 or more people (RUCA code = 1) and rural as anything not urban (RUCA codes 2-10). (Health Resources and Services Administration, n.d.). We conducted Bayesian group comparisons to assess the likelihood that each model performs differently by group.

Calibration and Overall Performance. To further evaluate our model's predictions, we used our inner resampling procedure (1 repeat of 5-fold cross validation) on the full data set to select and fit a single best model configuration for the full model. The final configuration selected represents the most reliable and robust configuration for deployment.

Calibration is an indicator of how well a model's predicted probabilities correspond to the true observed outcomes. For example, a well-calibrated model that assigns a 30% lapse risk prediction should observe lapses in approximately 30% of such cases. We used Platt scaling to calibrate our final full model's raw probabilities from the 5 held-out folds (Platt, 1999). We provided a calibration plot of these raw and calibrated probabilities. To characterize overall performance we reported Brier scores for the raw and calibrated probabilities. Brier scores are the mean squared difference between the predicted probabilities and observed outcome and range from 0 (perfect accuracy) to 1 (perfect inaccuracy). We also provided histograms of risk probability distributions by true lapse outcome.

Feature Importance

Feature importance values provide insight into the features that have the most influence on the model's predictions. For every prediction, we can extract feature importance values providing

actionable insight into intervenable targets for lapse risk (i.e., for a specific individal at a specific moment). We used the same single 5-fold cross-validation procedure (see Calibration and Overall Performance section) to calculate raw Shapley values for each observation in held-out data (Lundberg & Lee, 2017). The magnitude of the raw Shapley value indicates how much the feature score for that observation adjusted the prediction (in log-odds units) relative to the mean prediction across all observations. Positive Shapley values indicate that the feature score increased the prediction for that observation and negative values indicate that the feature score decreased the prediction. In other words, higher Shapley values suggest the feature increases lapse risk and lower values suggest the feature decreases lapse risk. Shapley values are inherently additive. For any observation, Shapley values can be summed to create a total adjustment score for the predicted value. We created feature categories by collapsing features that differed only by scoring epoch and calculation into a single feature category (e.g., Raw features for most recent urge rating, and min, max, and mean urge rating in the 168 hours preceding the start of the prediction window were combined to form a “Urge, Raw” feature category). Raw and difference features were not combined as their relationship to risk were not uniform. We plotted individual Shapley values and feature categories as partial dependence plots to illustrate these feature-risk relationships.

Feature importance values can also be aggregated across all participants and all observations to provide a relative rank ordering of the most important features. We calculated overall feature importance in two ways. First, we used a traditional approach in which we calculated the mean absolute Shapley value for each feature category across all observations. This

approach summarizes overall feature importance by averaging the magnitude of each feature's contribution. However, it can be skewed toward features that exhibit infrequent but very large Shapley values, potentially overstating the importance of features that are strongly associated with the outcome, but only come up in a small subset of observations. The second way we calculated feature importance was by calculating the proportion of observations in which each feature category had the highest Shapley value. This approach summarizes how frequently a feature category is influential across observations (i.e., considering both magnitude and prevalence). We provided a plot of the relative ranking of feature categories by their overall feature importance using these two methods.

Results

Participants

We enrolled 330 participants across 47 states in the United States from April 2021 through December 2024. To be considered enrolled in the study, participants were required to be eligible, consent, and complete the first personalized monthly survey where geolocation context were provided.[3] Of these 330, 14 were excluded for not providing at least 28 days of EMA data with acceptable adherence[4], 7 participants were excluded due to evidence of careless responding on daily surveys, and 13 participants were excluded due to having insufficient context data for geolocation points (fewer than two contextualized locations). Our final analysis sample consisted

[3] 405 participants consented to participate in the study and provided at least one daily survey; however, 75 participants dropped out prior to completing the first personalized monthly survey and were therefore not considered fully enrolled in the study.

[4] If adherence fell below 33% (roughly 9 surveys in a 28-day period) for a prolonged period of time, participant study end dates were changed to match the start of this adherence drop.

of 296 participants. Participant demographic and OUD characteristics of our final sample is presented in Table 1.

Table 1: Demographic and Clinical Characteristics

	N	%
Age		
22-25	7	2.4
26-35	105	35.5
36-45	107	36.1
46-55	60	20.3
56-65	13	4.4
Over 65	4	1.4
Gender		
Man	158	53.4
Woman	131	44.3
Non-binary	4	1.4
Not listed above	2	0.7
Orientation		
Straight, that is, not gay or lesbian	237	80.1
Lesbian or gay	17	5.7
Bisexual	35	11.8
Not sure	1	0.3
Not listed above	4	1.4
Race and Ethnicity (select all that apply)		
American Indian/Alaskan Native	17	5.7
Asian	3	1.0
Black/African American	43	14.5
Native Hawaiian/Other Pacific Islander	3	1.0
White/Caucasian	235	79.4
Hispanic, Latino, or Spanish origin	26	8.8
Not listed above	5	1.7
Geographic Location		
rural	62	20.9
urban	274	92.6
Education		
[illegible] teek	[illegible]	[illegible]
N = 296		

Adherence, Features, and Labels

Mean days on study across participants was 278.32 days (range 30-392 days). 77.7% of participants (230/296) remained on study for at least six months. Daily survey adherence was high. On average participants completed 77.8% of the daily survey prompts (range 40-100%). Participants provided, on average, 263.50 daily geolocation points (range 2.17-871.34).

The average percentage of missing values across daily survey features was low (mean=0.71%, range 0-1.10%). The average percentage of missing values across geolocation features was higher (mean=12.0%, range 0-15.6%). Across participants we generated a total of 82,156 day-level labels. Thirty-six percent of participants (106/296) reported an opioid lapse while on study (mean=4.93, range 0-117). This resulted in 1.8% of the labels positive for lapse (1,458/82,156 labels). We stratified the data on a variable of whether someone lapsed on study to ensure our imbalanced outcome was evenly split (and roughly equivalent to the sample prevalance) across cross-validation folds.

auROC and Model Comparisons

The median posterior probability for the best full model was 0.94, with narrow 95% CI ([0.93, 0.95]). To evaluate the incremental predictive value of our dynamic sensing features, we compared the performance of our full model to the performance of three ablated models: 0.63 (ablated daily survey features), 0.94 (ablated geolocation features), and 0.55 (ablated daily survey and geolocation features). There was strong evidence that the full model's performance was better compared to models that ablated the daily survey features and the daily survey and geolocation

features combined (probabilities = 1.00). There was no evidence that the full model performed better than a model with geolocation features ablated (probability = 0.34). The median difference in auROC, 95% Bayesian CI, and posterior probability that that the auROC difference was greater than 0 for all ablation contrasts are presented in Table 2.

Table 2: Median difference in auROC, 95% Bayesian credible interval (CI), and posterior probability that that the auROC difference was larger than 0 for all ablated model contrasts.

Contrast	Median auROC Difference	Bayesian CI	Probability
full vs. ablated daily survey	0.314	[0.29, 0.338]	1.000
full vs. ablated geolocation	−0.002	[−0.009, 0.005]	0.342
full vs. ablated daily survey and geolocation	0.389	[0.363, 0.414]	1.000
Median auROC differences greater than 0 indicate the full model, on average, performed better than the ablated model (i.e., full vs. ablated daily survey, full vs. ablated geolocation, full vs. ablated daily survey and geolocation). Bayesian CI represents the range of values where there is a 95% probability that the true auROC difference lies within that range. Probability indicates the posterior probability that this difference is greater than 0 (i.e., the full model is performing better).			

Our fairness subgroup comparisons for the full model revealed strong evidence that performance differed by gender, income, race and ethnicity, sexual orientation, and geographic location (probabilities > 0.97). Notably, our model performed better for individuals with an annual income of less than $25,000 compared to individuals with an annual income greater than $25,000, thus favoring the disadvantaged group. While differences in performance estimates exist, this model is likely still clinically meaningful across subgroups, as all subgroups yielded median auROCs between 0.91 - 0.95 (Figure 1). A table of fairness subgroup comparisons is available in the supplement.

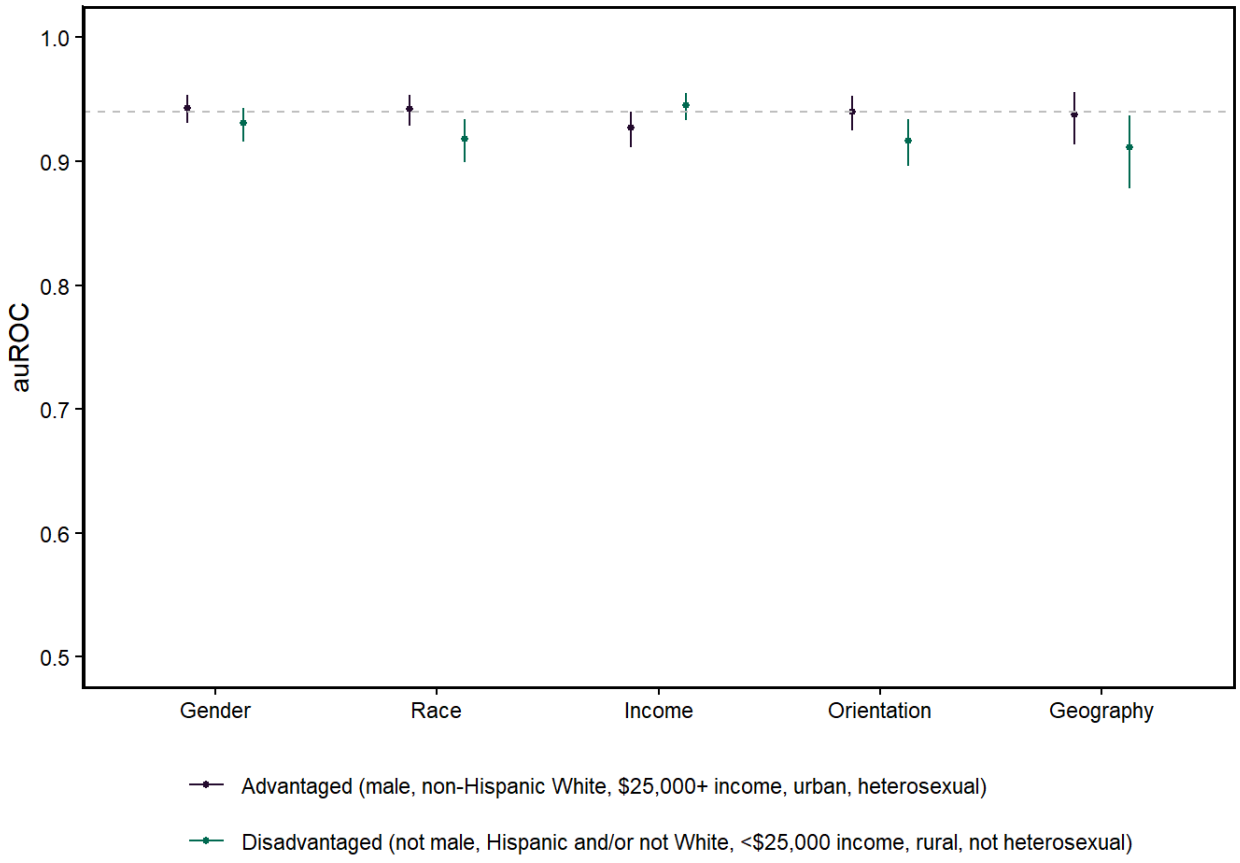

Figure 1: Posterior probabilities for area under the receiver operating curve (auROC) by demographic subgroup. auROC ranges from .5 (chance performance) to 1 (perfect performance). Subgroups advantaged in access to substance use treatment and outcomes (male, non-Hispanic White, greater than $25,000 annual income, urban location, and heterosexual) are depicted in dark purple. Subgroups disadvantaged in access to substance use treatment and outcomes (not male, Hispanic and/or not White, less than $25,000 annual income, rural location, and not heterosexual) are depicted in green. Overall model performance across groups is depicted as the dashed grey line.

Calibration and Overall Performance

The best full model configuration selected with 1 repeat of 5-fold cross-validation used an xgboost statistical algorithm and up-sampled the minority class.[5] This model produced generally well-calibrated raw probabilities (brier score = 0.014; left panel of Figure 2). We attempted to improve calibration with Platt scaling and results were comparable (brier score = 0.014). Histograms of raw risk probability distributions separately by true lapse outcome are presented in the right panel of Figure 2.

[5]The best model configuration used 1:5 upsampling of the majority class and the following hyperparameter values: learning rate = .1, tree depth = 3, mtry = 50.

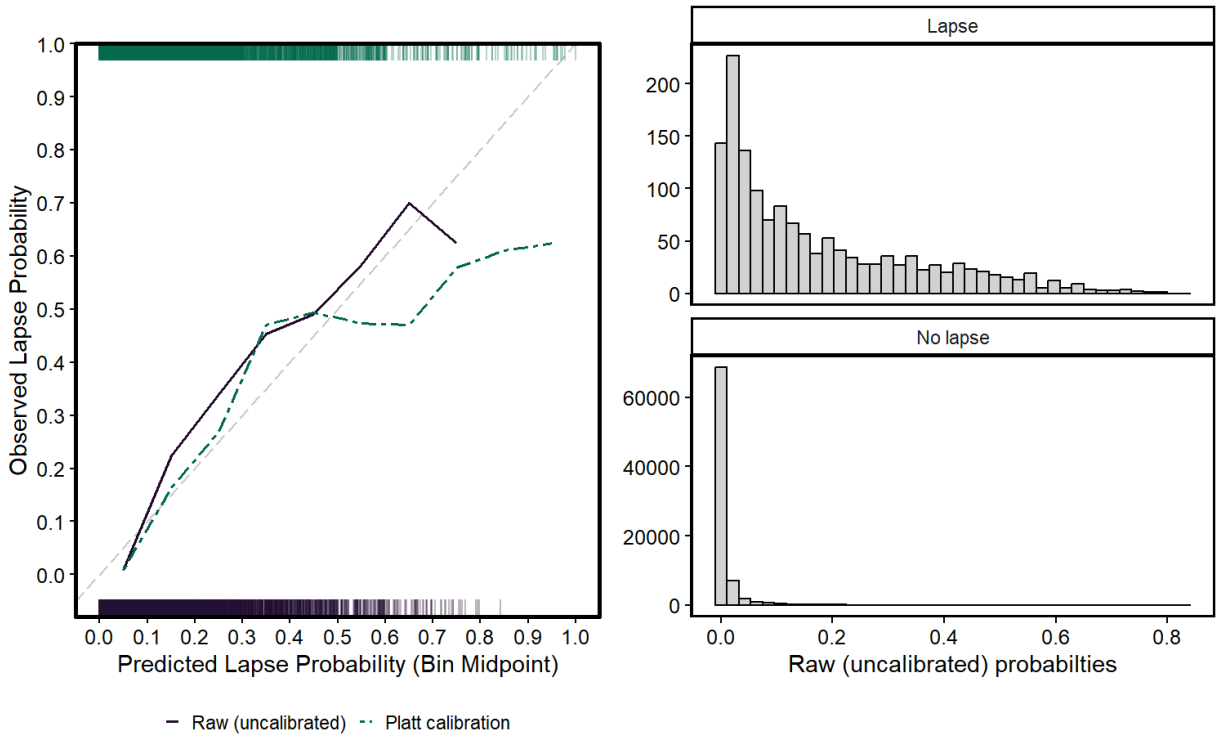

Figure 2: The left panel presents a calibration plot of raw and Platt scaled risk probabilities. Predicted probabilities (x-axis) are binned into deciles. Observed lapse probability (y-axis) represents the proportion of actual lapses observed in each bin. The dashed diagonal represents perfect calibration. Points below the line indicate overestimation and points above the line indicate underestimation. Raw probabilities are depicted as the dark purple line. Platt calibrated probabilities are depicted as the green dashed line. The rug plots along the x-axes depict the observation frequency in each bin for the raw (bottom) and Platt scaled (top) probabilities. The right panel presents histograms of raw (uncalibrated) risk probability distributions separately by true lapse outcome.

Feature Importance

We aggregated individual Shapley values across all participants and all observations to provide a relative rank ordering of the most important features. A traditional mean absolute Shapley value approach revealed the most important feature (i.e., largest magnitude of change in log-odds) was differences in past opioid use. These feature was nearly two-fold more important than the next most important feature, raw urge ratings. Figure 3 displays the top 30 (out of 186) feature categories, as defined by the 30 largest absolute mean shapley values. In the supplement, we provide full figures of all possible feature categories.

A second approach where we calculated the proportion of observations in which each feature category had the highest Shapley value (i.e., largest magnitude and prevalence), revealed the most important feature to be raw urge ratings, followed by raw abstinence confidence ratings. Differences in past opioid use was the third most important feature by this method, suggesting that while it has a strong relationship with lapse risk, it is not as prevalent across observations as other features. Notably, in this approach several additional comparably important features emerged, including differences in past pleasant experiences, differences in past stimulant use, and differences in time spent in helpful locations.

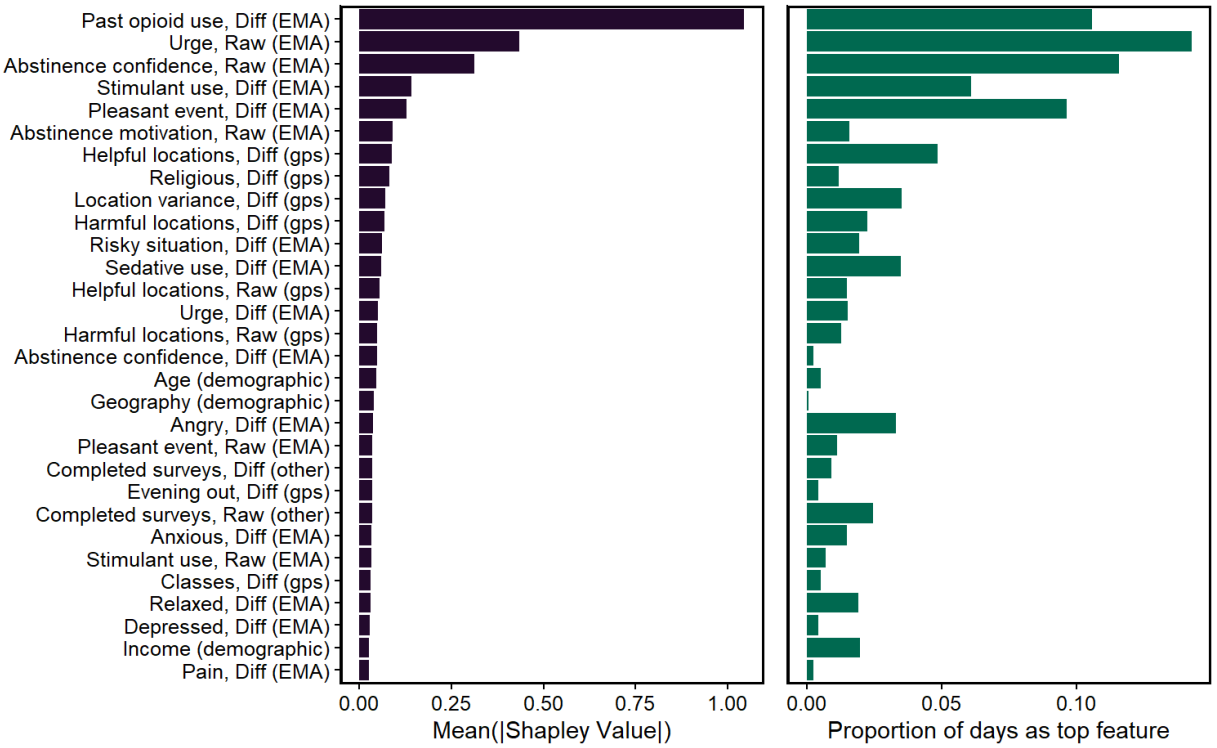

Figure 3

To understand the direction in which these feature categories adjusted the lapse risk predictions, we plotted Shapley-feature value partial dependence plots for the top 30 feature categories (Figure 4). Differenes in past opioid use lowered lapse risk probability for participants who did not lapse on study (i.e., a feature score of 0). Lapse probability dramatically increased when participants had increased reported opioid use over the past 24 and 168 hours. Unsurprisingly, raw urge and abstinence confidence ratings had linear relationships to lapse risk, with higher urge ratings increasing lapse risk probability and higher abstinence confidence decreasing lapse risk probability. There was a more nuanced finding for stimulant use and opioid lapse risk probability.

Raw rate counts of stimulant use increased lapse risk; however, participants were at highest risk when recent stimulant use decreased (i.e., Stimulant Use, Diff).

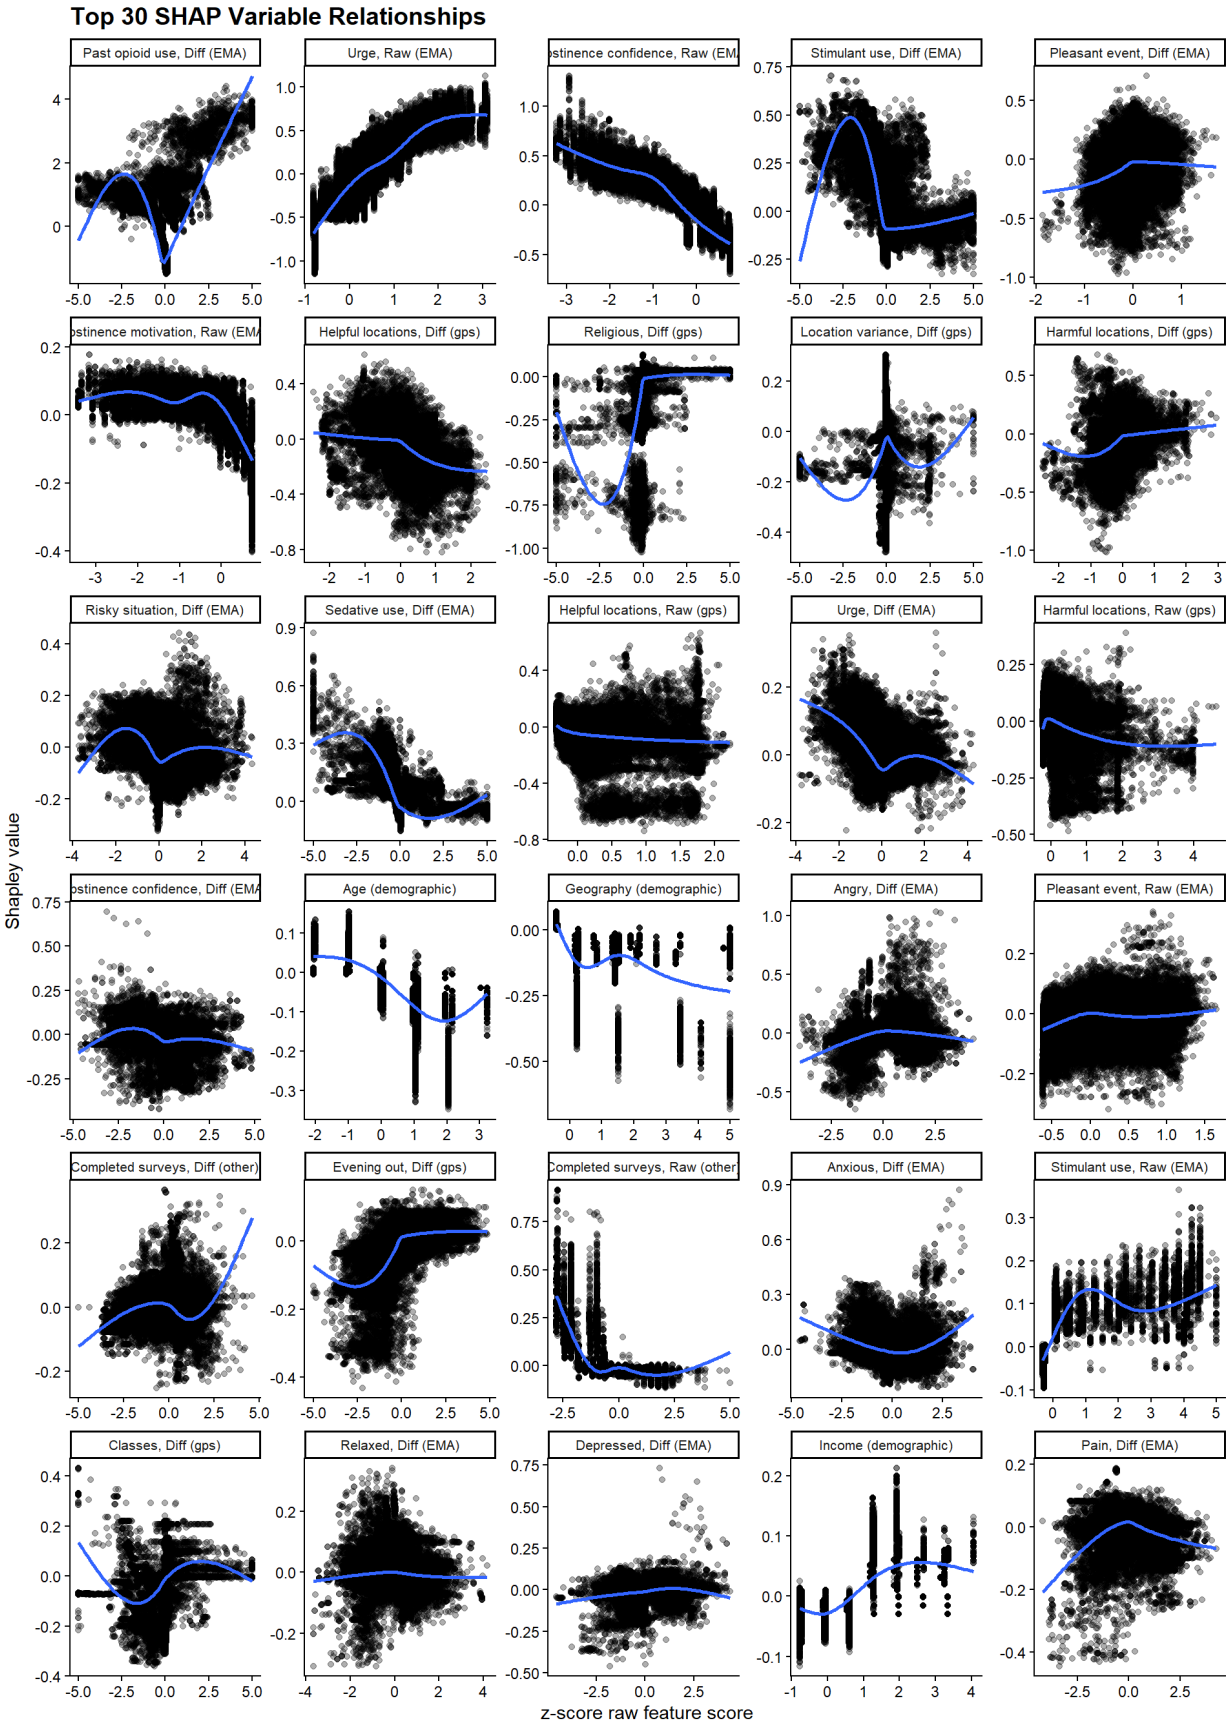

Figure 4: Feature importance partial dependence plots.

Discussion

- Discuss contrast between top global features and “actual” important features (i.e., what comes up day to day as being important). What does this mean for how we traditionally conceptualize feature importance?
- day of prediction window not important for oud but important for aud (add supplemental figure of lapses by day of week and hour?)
- When discussing fairness - we don’t look at age. People over 65 is an important group to look at but they were not in our sample
- Acknowledge fairness differences still exist. Research should look into additional risk factors that may better predict risk of lapse in minority populations.
- Decision not include misclassification cost analyses: acknowledge that discrimination and calibration don’t equate to clinical utilility. Decision analysis tools, such as net benefit curves could be used to quantify clinical benefit at relevant probability thresholds. Our model is designed to not be used to make decisions about whether or not to treat but to provide model feedback to individuals each day to help them monitor their risk for lapse. Therefore the decision has already been made to always treat. Should these models be used to inform when to deploy more cost-intensive interventions (e.g., communicating risk level to a therapist to initiate contact) these analyses will be important next steps.
- Although, classification and clinical utility are two other important domains for evaluating a model intended to make or inform a decision (e.g., whether to send a support message to an individual). In applications of our model there would be no decision to be made. Every day an individual would receive a message regardless of their lapse risk, equivalent to a treat all

condition. Therefore, we have constrained our evaluation to focus on the probability estimates for evaluating performance.

- Optimal thresholds depend on the setting, outcome prevalence, and the relative costs of misclassification, and should be addressed separately from discrimination evaluation (e.g., with a decision curve analysis).
- Discussion about GPS data. No evidence that geolocation offered incremental predictive value. Hard to get. With better sensing features could be more important. 12% of gps features were imputed as median. Low leverage for predictions.
- If all you care about is prediction, you don't need gps. If you desire a more nuanced understanding of risk across a more diverse array of risk features than gps is helpful.
- GPS may be better if we did not rely on only known locations but could passively contextualize locations. May be better with more robust sensing methods.
- Overall, GPS didn't hurt predictions but also didn't add any value above and beyond daily surveys. Appears to overlap with EMA.
- Daily surveys remain to be a feasible and predictive signal even when measured over a year.
- Participants provided data for up to 12 months. This extended window of recovery (12 months vs. 3 months) is critical for evaluating the value of an algorithm intended for ongoing continuing care support and understanding how lapse risk evolves as people progress in their recovery (i.e., past 3 months). Unfortunately, our ability to address explanatory questions about the time course of lapse risk, and how individuals might cluster on different recovery trajectories is limited with traditional machine learning methods. These methods do not capture the repeated nature of sensing data. Each lapse prediction is treated as a new independent observation. We can account for the repeated observations in our sensing data by engineering

features that capture individual changes over time to produce unbiased and precise estimates of predictive performance. However, more traditional time series models are better suited for understanding the temporal dynamics of lapse risk over long periods of recovery.

- The longer duration of participation also provides the opportunity to experiment with lagged prediction windows further than two weeks into the future (i.e., 1 month). However, in a machine learning framework, different models must be built to predict lapses at varying times in the future. For example, a recovery monitoring support system that detects both immediate lapse risk (i.e., in the next 24 hours) and future lapse risk (i.e, in the next 2 weeks and in the next month) would need three models. This approach is time consuming, cumbersome, and still only provides coarse understanding of time-course. Therefore, this study will only predict immediate lapse risk.

References

American Psychiatric Association. (2022). *Diagnostic and statistical manual of mental disorders : DSM-5-TR*. 5th edition, text revision. Washington, DC : American Psychiatric Association Publishing, 2022.

Center for High Throughput Computing. (2006). *Center for high throughput computing*. Center for High Throughput Computing. https://doi.org/10.21231/GNT1-HW21

Collins, G. S., Moons, K. G. M., Dhiman, P., Riley, R. D., Beam, A. L., Van Calster, B., Ghassemi, M., Liu, X., Reitsma, J. B., Van Smeden, M., Boulesteix, A.-L., Camaradou, J. C., Celi, L. A., Denaxas, S., Denniston, A. K., Glocker, B., Golub, R. M., Harvey, H., Heinze, G., … Logullo, P. (2024). TRIPOD+AI statement: Updated guidance for reporting clinical prediction

models that use regression or machine learning methods. *BMJ*, *385*, e078378. https://doi.org/10.1136/bmj-2023-078378

Economic Research Service US Department of Agriculture. (n.d.). *Rural-Urban Commuting Area Codes - Users' Guide*. https://www.ers.usda.gov/data-products/rural-urban-commuting-area-codes/users-guide.

Goodrich, B., Gabry, J., Ali, I., & Brilleman, S. (2023). *Rstanarm: Bayesian Applied Regression Modeling via Stan*.

Greenfield, S. F., Brooks, A. J., Gordon, S. M., Green, C. A., Kropp, F., McHugh, R. K., Lincoln, M., Hien, D., & Miele, G. M. (2007). Substance abuse treatment entry, retention, and outcome in women: A review of the literature. *Drug and Alcohol Dependence*, *86*(1), 1–21. https://doi.org/10.1016/j.drugalcdep.2006.05.012

Gustafson, D. H., McTavish, F. M., Chih, M.-Y., Atwood, A. K., Johnson, R. A., Boyle, M. G., Levy, M. S., Driscoll, H., Chisholm, S. M., Dillenburg, L., Isham, A., & Shah, D. (2014). A smartphone application to support recovery from alcoholism: A randomized clinical trial. *JAMA Psychiatry*, *71*(5), 566–572. https://doi.org/10.1001/jamapsychiatry.2013.4642

Health Resources and Services Administration. (n.d.). *How We Define Rural*. https://www.hrsa.gov/rural-health/about-us/what-is-rural.

Kilaru, A. S., Xiong, A., Lowenstein, M., Meisel, Z. F., Perrone, J., Khatri, U., Mitra, N., & Delgado, M. K. (2020). Incidence of Treatment for Opioid Use Disorder Following Nonfatal Overdose in Commercially Insured Patients. *JAMA Network Open*, *3*(5), e205852. https://doi.org/10.1001/jamanetworkopen.2020.5852

Kuhn, M. (2022). *Tidyposterior: Bayesian Analysis to Compare Models using Resampling Statistics*.

Kuhn, M., & Wickham, H. (2020). *Tidymodels: A collection of packages for modeling and machine learning using tidyverse principles*.

Lofaro, R. J., Bohler, R. M., Spurgeon, R., & Mase, W. A. (2025). Rural and urban differences in treatment on demand for substance use treatment involving medications for opioid use disorder. *The Journal of Rural Health*, *41*(3), e70076. https://doi.org/10.1111/jrh.70076

Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Proceedings of the 31st International Conference on Neural Information Processing Systems, NIPS'17*, 4768–4777.

Martin, C. E., Parlier-Ahmad, A. B., Beck, L., Scialli, A., & Terplan, M. (2022). Need for and Receipt of Substance Use Disorder Treatment Among Adults, by Gender, in the United States. *Public Health Reports*, *137*(5), 955–963. https://doi.org/10.1177/00333549211041554

Moshontz, H., Colmenares, A. J., Fronk, G. E., Sant'Ana, S. J., Wyant, K., Wanta, S. E., Maus, A., Jr, D. H. G., Shah, D., & Curtin, J. J. (2021). Prospective Prediction of Lapses in Opioid Use Disorder: Protocol for a Personal Sensing Study. *JMIR Research Protocols*, *10*(12), e29563. https://doi.org/10.2196/29563

Olfson, M., Mauro, C., Wall, M. M., Choi, C. J., Barry, C. L., & Mojtabai, R. (2022). Healthcare coverage and service access for low-income adults with substance use disorders. *Journal of Substance Abuse Treatment*, *137*, 108710. https://doi.org/10.1016/j.jsat.2021.108710

Pinedo, M. (2019). A current re-examination of racial/ethnic disparities in the use of substance abuse treatment: Do disparities persist? *Drug and Alcohol Dependence*, *202*, 162–167. https://doi.org/10.1016/j.drugalcdep.2019.05.017

Platt, J. (1999). Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in large margin classifiers* (pp. 61–74). MIT Press.

Van Calster, B., Collins, G. S., Vickers, A. J., Wynants, L., Kerr, K. F., Barreñada, L., Varoquaux, G., Singh, K., Moons, K. G., Hernandez-Boussard, T., Timmerman, D., McLernon, D. J., Van Smeden, M., & Steyerberg, E. W. (2025). Evaluation of performance measures in predictive artificial intelligence models to support medical decisions: Overview and guidance. *The Lancet Digital Health*, 100916. https://doi.org/10.1016/j.landig.2025.100916

Vandewiele, G., Dehaene, I., Kovács, G., Sterckx, L., Janssens, O., Ongenae, F., De Backere, F., De Turck, F., Roelens, K., Decruyenaere, J., Van Hoecke, S., & Demeester, T. (2021). Overly optimistic prediction results on imbalanced data: A case study of flaws and benefits when applying over-sampling. *Artificial Intelligence in Medicine*, *111*, 101987. https://doi.org/10.1016/j.artmed.2020.101987

Xin, Y., Schwarting, C. M., Wasef, M. R., & Davis, A. K. (2023). Exploring the intersectionality of stigma and substance use help-seeking behaviours among lesbian, gay, bisexual, transgender, queer, questioning or otherwise gender or sexuality minority (LGBTQ+) individuals in the United States: A scoping review. *Global Public Health*, *18*(1), 2277854. https://doi.org/10.1080/17441692.2023.2277854